

Machine Learning Task

Smart Targeting for Bank Term Deposits

Instructions for the candidate: Choose **either 1 or 2** and complete it in **30 minutes**. You may use Python/R. Keep it simple and readable.

Dataset: Kaggle — Bank Marketing (code tab):

<https://www.kaggle.com/datasets/henriqueyamahata/bank-marketing/data>

Option 1 — Single-Campaign Ranking (hands-on, coding)

Goal. Train a quick binary model that ranks customers for a term-deposit campaign.

Timebox. 30 minutes. **Lang.** Python or R. **Keep it simple.**

Data + Split

- Load the Kaggle Bank Marketing dataset (link above).
- **Exclude leakage:** drop **duration** from features.
- Train/test split: stratified by target **y**, with a fixed random seed (e.g., 42).

Minimal Preprocessing

- One-hot encode categoricals; keep numerics as is (standardize only if your model needs it).
- Optional: handle rare categories by “other” or let one-hot handle naturally.

Model

- Train one fast classifier (e.g., Logistic Regression with `class_weight="balanced"`, or a small tree/GBM).
- Use the train set only (no peeking).

Scoring + Metrics

- Score the **test** set → predicted probability $p = P(y=1)$.
- Compute:
 - **ROC-AUC** (and/or PR-AUC; ROC-AUC required).
 - **Top-k precision** where $k = \text{round}(0.10 * \text{test_size})$.
- Build a **Top-15 table** sorted by p desc with columns:
`row_id (or customer_id), predicted_prob.`

What to Print (exact)

Print these lines (replace ALL CAPS with your values):

- `AUC_ROC: XX.XXX`
- `TopK_Precision@10%: XX.XXX`

Then display:

- **Top 15 table:** `row_id, predicted_prob` (rounded to 3–4 decimals).

Short Notes (final cell, 3–5 lines total)

- **Which features mattered:** list 2–3 (e.g., top coefficients or importances).
- **How we'd deploy daily:** batch scoring on latest data, rank by p , contact top N given capacity; monitor AUC, conversion rate, and drift.

Option 2 — Multi-Campaign & Channel Policy (hands-on, coding)

Goal. Learn a policy to pick (`campaign_intensity`, `channel`) per customer to maximize conversions.

Timebox. 30 minutes. **Lang.** Python or R. **Keep it fast.**

Tip. Don't use `duration` (leakage).

Setup

- Load data; stratified train/test split on `y`, fixed seed.
- Features = pre-contact info only (exclude `duration`). One-hot categoricals.
- Define **actions** from data:
 - `campaign_intensity` $\in \{1, 2+\}$ derived from `campaign`.
 - `channel` $\in \{\text{cellular}, \text{telephone}\}$ from `contact`.
 - Actions: **A1** (1, cellular), **A2** (1, telephone), **A3** (2+, cellular), **A4** (2+, telephone).
- Reward: `y == "yes"` \rightarrow 1 else 0.

Modeling (choose one)

- **Per-action models:** train 4 binary models, each on rows where that action occurred; at inference, score all 4 and pick $\text{argmax } P(y=1 \mid \text{action})$.
OR
- **Shared policy:** train a multiclass policy to choose the action; also train a single binary outcome model for $P(y=1)$ to report probabilities/ties.

Offline Eval (implement end-to-end)

- **Diagnostics:** ROC-AUC per binary model (or for your outcome model). Print class balance and note if you used class weights.
- **Policy value (action-match):** apply policy to test, keep rows where chosen action equals logged action, compute `conversions / matched_rows`. Also compute two baselines: random action (4-way) and a fixed policy (e.g., always (1, cellular)).
- **Top-k lift:** define a **policy score** per user (e.g., max predicted $P(y=1|a)$). With `k = round(0.10 * test_size)`, compute Top-k precision for policy vs random; report **lift** = `policy / random`.
- **Recommendation table (10 rows):** columns = `row_id`, `chosen_campaign_intensity`, `chosen_channel`, `policy_score`, `reason`. Show top 10 by `policy_score`.
- **Feature importance:** quick print (coefficients or importances).

What to Print (exact)

- `AUC_per_action: {A1: XX.XXX, A2: XX.XXX, A3: XX.XXX, A4: XX.XXX}`
- `PolicyValue_Match: XX.XXX`
- `Baseline_Random_Match: XX.XXX`
- `Baseline_Fixed_Match: XX.XXX`
- `TopK_Precision (policy): XX.XXX`
- `TopK_Precision (random): XX.XXX`
- `Lift: XX.XXX`

Then display:

- The **10-row** recommendation table.
- A short top-features printout (per model or overall).